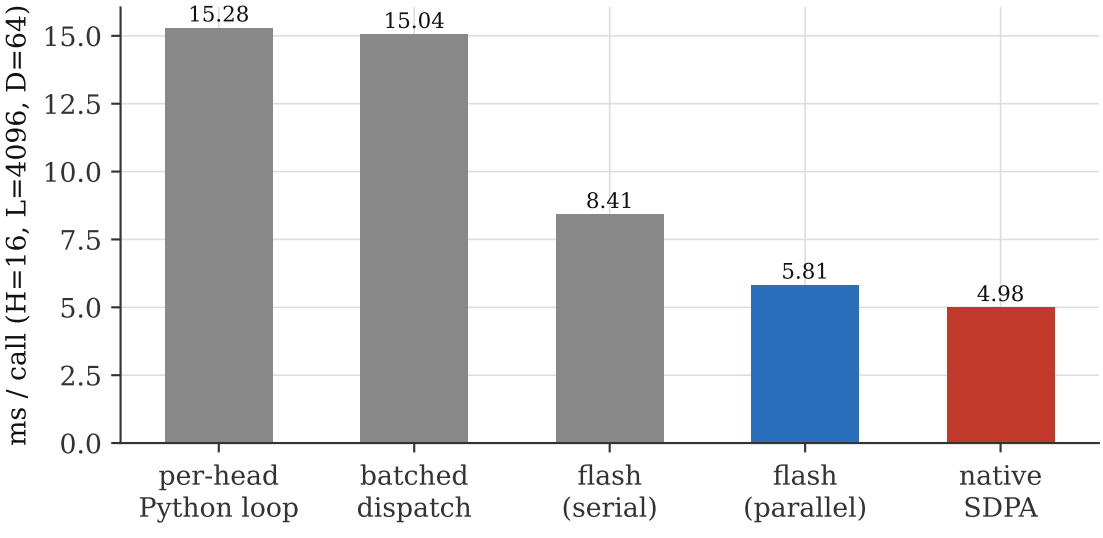
Attention kernel progression toward native
ms / call (H=16, L=4096, D=64)
0.0
2.5
5.0
7.5
10.0
12.5
15.0
15.28
15.04
8.41
5.81
4.98
per-head
Python loop
batched
dispatch
flash
(serial)
flash
(parallel)
native
SDPA